


Sakkara

A hierarchical cluster topology group scheduler




Topology-Aware Workload Placement

- Placing a workload impacts its performance
- Sakkara
 - solves the placement problem for all pods at once accounting for
 - resource requirements
 - topology constraints
 - place each pod accordingly



topology specification



```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: kind-topology-configmap
5    namespace: default
6    labels: {
7      "sakkara.topology":
8    }
9  data:
10   name: "kind-tree"
11   resource-names: |
12     [ "cpu", "memory", "nvidia.com/gpu" ]
13   level-names: |
14     [ "rack", "node" ]
15   tree: |
16     {
17       "rack-0": {
18         "cluster1-control-plane": {}
19       },
20       "rack-1": {
21         "cluster1-worker": {},
22         "cluster1-worker2": {}
23       },
24       "rack-2": {
25         "cluster1-worker3": {},
26         "cluster1-worker4": {}
27       }
28     }
```

label used for filtering

ordered resource names

level names top down

tree hierarchical topology

Sakkara: Highlights

- Scheduler plugin
 - Hierarchical cluster topology
 - configmap -> tree specs
 - Group (gang) scheduling
 - configmap -> group specs, placement results
 - or PodGroup
 - Logical application topology generation (pods ranking)
 - Uses chic-sched as core group placer (solver)
- Supports
 - Dynamic tree cluster topology
 - Multiple resources
 - Static, homogeneous groups
 - Group placement constraints
 - multi-level, multi-type (spread, pack, partition, range, factor), hard/soft
 - Group preemption (all or none)
 - Weighted nodes

group specification

label used
for filtering

group name,
size, priority

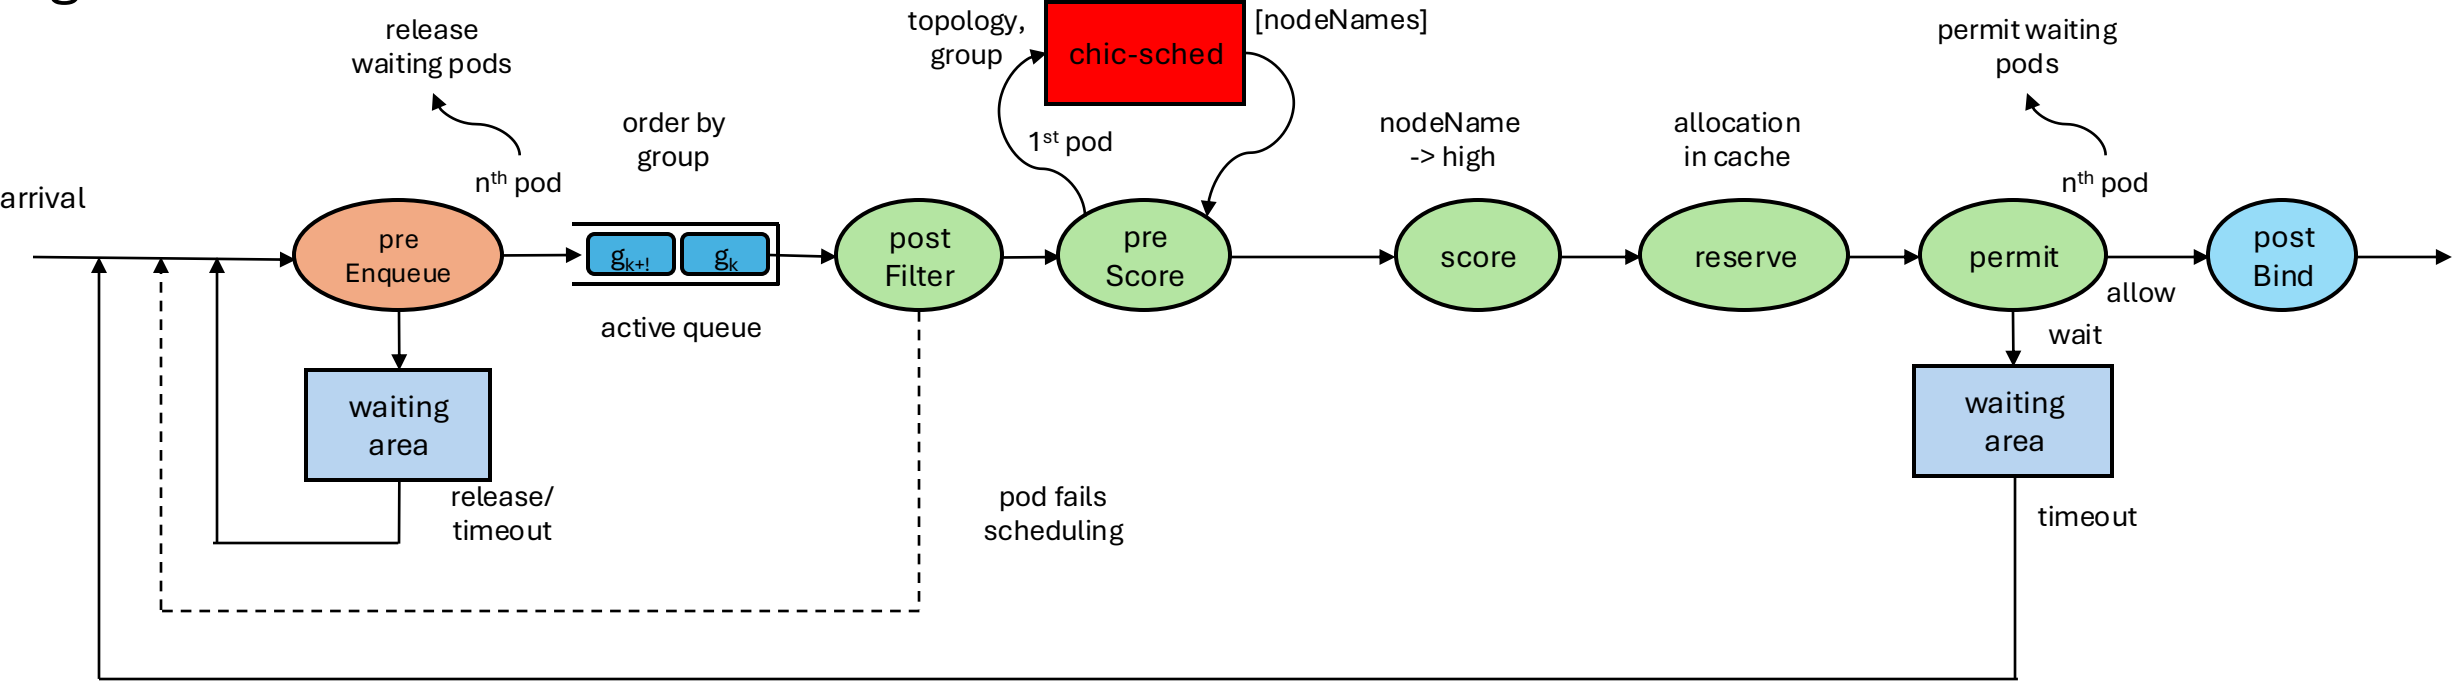
group level
constraints

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: group-a-0
5    namespace: default
6    labels: {
7      "sakkara.group":
8    }
9  data:
10   "name": "group-a-0"
11   "size": "6"
12   "priority": "0"
13   "constraints": |
14     {
15       "rack": {
16         "type": "spread"
17       },
18       "node": {
19         "type": "pack"
20       }
21     }
22
```

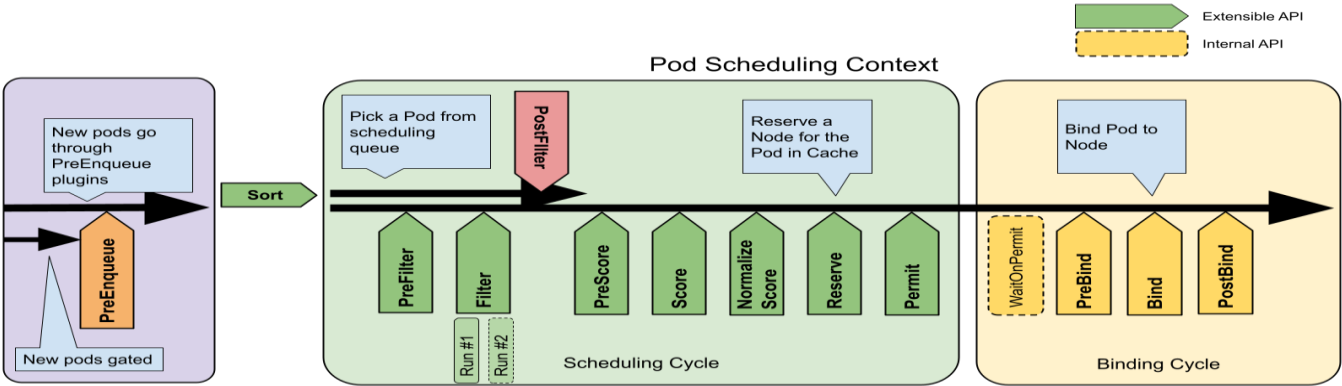
job specification

```
23  kind: Deployment
24  apiVersion: apps/v1
25  metadata:
26    name: deploy-a-0
27    namespace: default
28  spec:
29    replicas: 6
30    selector:
31      matchLabels:
32        app: group-a-0
33    template:
34      metadata:
35        labels:
36          app: group-a-0
37          sakkara.group.name: "group-a-0"
38    spec:
39      schedulerName: sakkara
40      imagePullSecrets:
41        - name: regcred
42      containers:
43        - name: container-1
44          image: nginx
45          imagePullPolicy: IfNotPresent
46          ports:
47            - name: web
48              containerPort: 80
49              protocol: TCP
50          resources:
51            requests:
52              cpu: "0.2"
53              memory: "200Mi"
54              nvidia.com/gpu: "2"
55            limits:
56              nvidia.com/gpu: "2"
```

plugin state transitions



group
status



job scheduling results

```
I1129 17:20:35.919734 61198 solver.go:89] "Solve: " pTree=<
```

```
pTree:
root -> ( rack-1 -> ( cluster1-worker cluster1-worker2 ) rack-2 -> ( cluster1-worker3 cluster1-worker4 ) )
pNodes:
pNode: ID=root; level=2; cap=[4000 7921025024 32]; alloc=[1600 1468006400 12]; numClaimed=6
pNode: ID=rack-1; level=1; cap=[2000 3960512512 16]; alloc=[800 734003200 6]; numClaimed=3
pNode: ID=rack-2; level=1; cap=[2000 3960512512 16]; alloc=[800 734003200 6]; numClaimed=3
pNode: ID=cluster1-worker; level=0; cap=[1000 1980256256 8]; alloc=[700 681574400 6]; numClaimed=3
pNode: ID=cluster1-worker2; level=0; cap=[1000 1980256256 8]; alloc=[100 52428800 0]; numClaimed=0
pNode: ID=cluster1-worker4; level=0; cap=[1000 1980256256 8]; alloc=[700 681574400 6]; numClaimed=3
pNode: ID=cluster1-worker3; level=0; cap=[1000 1980256256 8]; alloc=[100 52428800 0]; numClaimed=0
```


```
> pg=<
```

```
PG: ID=group-a-0; size=6; demand=[200 209715200 2]; lcs=[node rack]
lTree:
root -> ( rack-1 -> ( cluster1-worker ) rack-2 -> ( cluster1-worker4 ) )
lNodes:
lNode: ID=root; count=6; claimed=6
lNode: ID=rack-1; count=3; claimed=3
lNode: ID=rack-2; count=3; claimed=3
lNode: ID=cluster1-worker; count=3; claimed=3
lNode: ID=cluster1-worker4; count=3; claimed=3
```

```
> lcs=["rack","node"]
```

```
I1129 17:20:35.919748 61198 solver.go:92] "Solve: " levelConstraint="LC: ID=node; level=0; affinity=Pack; isHard=false; "
I1129 17:20:35.919752 61198 solver.go:92] "Solve: " levelConstraint="LC: ID=rack; level=1; affinity=Spread; isHard=false; "
```

physical tree:
resource allocation



logical tree:
group placement

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
deploy-a-0-759894d8c8-2h2cr	1/1	Running	0	3m51s	10.244.4.4	cluster1-worker
deploy-a-0-759894d8c8-4h9vv	1/1	Running	0	3m51s	10.244.2.3	cluster1-worker4
deploy-a-0-759894d8c8-62k2x	1/1	Running	0	3m51s	10.244.4.2	cluster1-worker
deploy-a-0-759894d8c8-bwjxh	1/1	Running	0	3m51s	10.244.2.2	cluster1-worker4
deploy-a-0-759894d8c8-j9bsx	1/1	Running	0	3m51s	10.244.4.3	cluster1-worker
deploy-a-0-759894d8c8-xxk6v	1/1	Running	0	3m51s	10.244.2.4	cluster1-worker4

group configmap

```
apiVersion: v1
data:
  constraints: |
    {
      "rack": {
        "type": "spread"
      },
      "node": {
        "type": "pack"
      }
    }
  name: group-a-0
  placement: '{"root":{"rack-1":{"cluster1-worker":{"deploy-a-0-759894d8c8-2h2cr":{},"deploy-a-0-759894d8c8-62k2x":{},"deploy-a-0-759894d8c8-j9bsx":{}}},"rack-2":{"cluster1-worker4":{"deploy-a-0-759894d8c8-4h9vv":{},"deploy-a-0-759894d8c8-bwjxh":{},"deploy-a-0-759894d8c8-xxk6v":{}}}}}'
  priority: "0"
  rank: '[(deploy-a-0-759894d8c8-2h2cr,0) (deploy-a-0-759894d8c8-62k2x,1) (deploy-a-0-759894d8c8-j9bsx,2) (deploy-a-0-759894d8c8-4h9vv,3) (deploy-a-0-759894d8c8-bwjxh,4) (deploy-a-0-759894d8c8-xxk6v,5)]'
  size: "6"
  status: Bound
kind: ConfigMap
metadata:
  creationTimestamp: "2023-11-29T22:20:35Z"
  labels:
    sakkara.group: ""
  name: group-a-0
  namespace: default
  resourceVersion: "1510"
  uid: c515880e-08c3-4747-8c2e-4a6bcb49071c
```

Pods
placement
rank

group
status

rank: order based on (hierarchical)
distance from master pod (rank 0)

pod labels

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2023-11-29T22:20:35Z"
  generateName: deploy-a-0-759894d8c8-
  labels:
    app: group-a-0
    pod-template-hash: 759894d8c8
    sakkara.group.name: group-a-0
    sakkara.member.rank: "0"
    sakkara.member.retries: "1"
    sakkara.member.status: Bound
  name: deploy-a-0-759894d8c8-2h2cr
  namespace: default
```

group name

sakkara-
generated
attributes

pod name

Sakkara: Algorithm

chic-sched:

A scheduler for HPC placement groups

<https://github.ibm.com/chic/chic-sched>

- Heuristics-based and topology-aware group placement algorithm
 - physical and logical trees (infrastructure-concealed, application topology)
 - multiple-level constraints (pack racks, spread servers)
 - tree traversal without retrials, suboptimal, fast, depends on heuristics
 - modeling and analysis driven heuristics design
- Dynamic (elastic) placement groups
 - optimal addition/deletion
 - provisional scheduling

```
kind: GroupPlacement
spec:
  group:
    name: MyApp
    size: 20
    type: bx2-16x64
  constraints:
    - level: rack
      affinity: pack
      soft: true
```

group size

- specified, fixed
- unspecified, estimated, dynamic


type

- homogeneous
- specifies resources requested

level: zone | room | rack | server

affinity: spread | pack


soft: true | false



Simulation Experiment I



Infrastructure:
1 zone
3 rooms/zone
8 racks/room
20 servers/rack

Placement Group:
120 VMs
1/8 server



Spread Rack Soft

Pack Server Soft

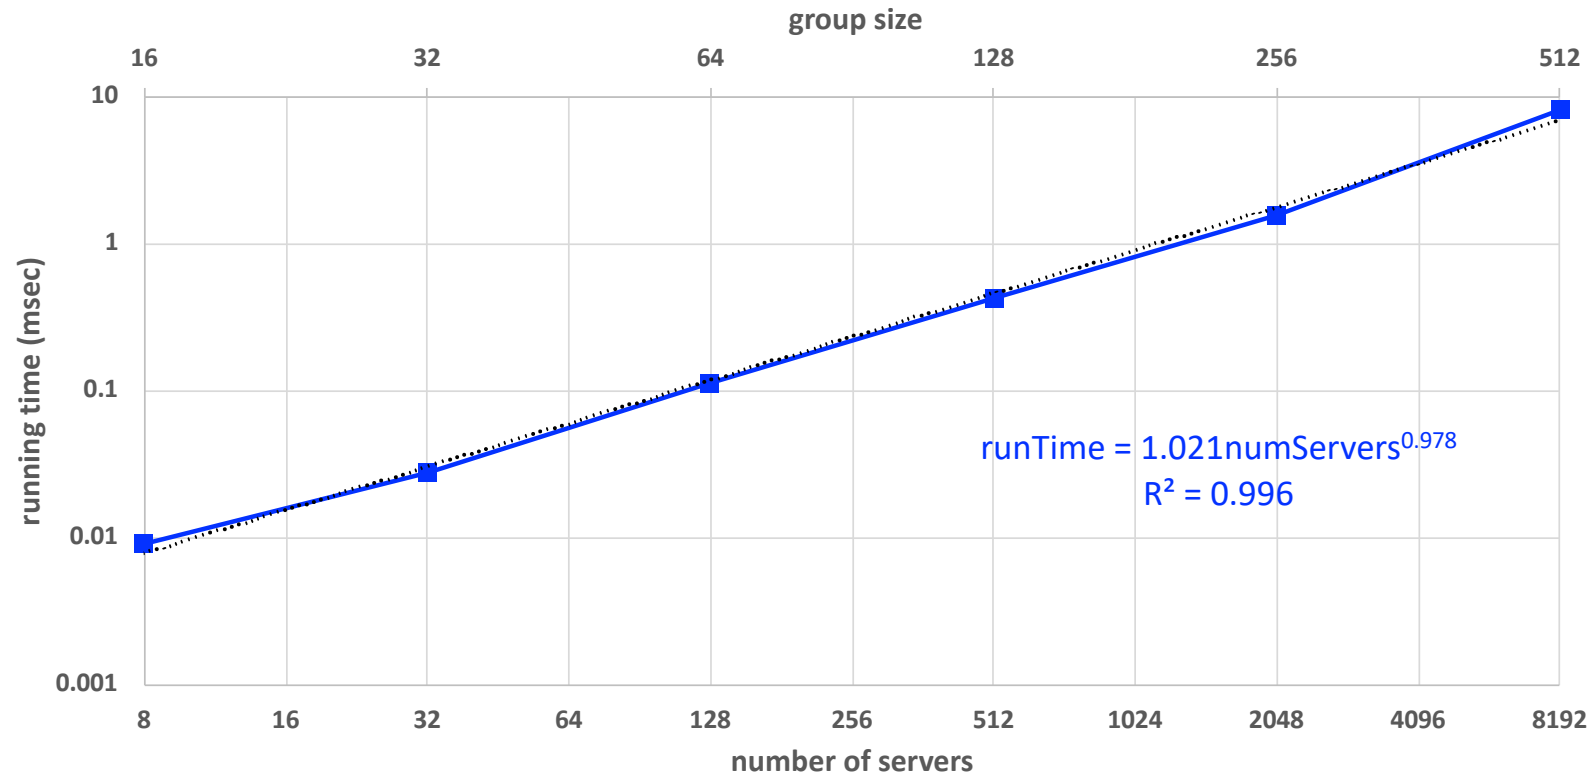


Placement algorithm

- Traverse physical tree in depth first order
- Solve placement subproblem when visiting node n
- Choice of heuristics
 - ordering of sibling nodes
 - determination of placement range
 - selection of best number to place at a node
- Return of logical tree, representing placement result
- Variations of algorithm
 - place a partially placed group
 - place dynamic group (size changes)

Performance

Scalability of group placement algorithm



Configuration

- 1 zone, x racks, 2x servers/rack
- group size 8x
- resources
 - CPU
 - capacity = [16]
 - demand = [2]
- load = 40%
- skewed allocation
 - $p_0=1/4$, $p_1=1/8$, $\text{CoV}=1.5$
- constraints
 - rack, spread, soft
 - server, pack, soft
- 10,000 samples per point